

APRIL 2023



“RAYS THE BAR”

HW Acceleration of Ray-Tracing on an FPGA System

ECE 382N – Adv Embedded Microcontroller Sys

Course Project Presentation


Ganesh Ram Koushik, (gk7734)

Shreyas Ravishankar, (sr48925)

The University of Texas at Austin

What is Ray Tracing?

- Computer Graphics technique for Rendering “Realistic” images.
- Simulates the interaction between light rays and the objects.
- In simple terms, it is a Ray Optics physics model.



High Level Overview


- Implemented SW code for Ray Tracing.
- Used Verilog and HLS to offload sub-tasks to the HW.
- Achieved **5.07x** speedup over baseline software running on Ultra96.

Algorithm Discussion

- The PPM Format
- The Ray Model
- The Camera & Viewport
- Ray-Object Interactions
- Anti-Aliasing
- Surfaces

The PPM Format

```
Sample.ppm X
Sample.ppm
1 p3
2 3 2
3 255
4 255 0 0
5 0 255 0
6 0 0 255
7 255 255 0
8 255 255 255
9 0 0 0
```




The Ray Model

$$\vec{v} = \vec{A} + \vec{B} * t$$


\vec{A} is the Ray Origin

\vec{B} is the Ray Direction



The Camera and the Viewport

- A scene is described using a global coordinate system.
- A ray is shot into every pixel of the viewport.
- Color of the pixel is determined by the object(s) that this ray “hits”.



Ray Object Interactions


$$(x - C_x)^2 + (y - C_y)^2 + (z - C_z)^2 = r^2$$

Since the Intersection Point can be represented as a Ray

$$(P_x(t) - C_x)^2 + (P_y(t) - C_y)^2 + (P_z(t) - C_z)^2 = r^2$$

Simplifying

$$t^2 * (B \cdot B) + 2t * ((B) \cdot (A - C)) + (A - C) \cdot (A - C) = r^2$$



Anti-Aliasing

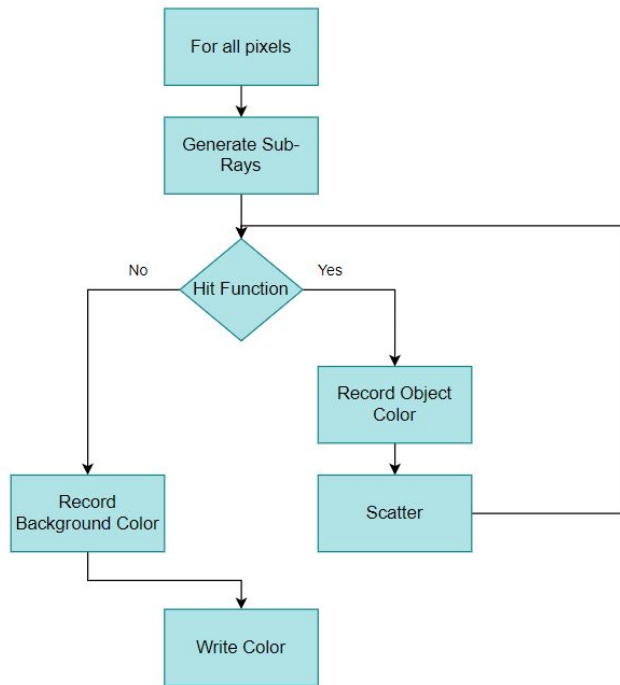
- Pixels at object borders contain multiple colors.
- Shooting one ray per pixel renders objects with jagged edges.
- Shoot multiple (~ 100) rays per pixel to smoothen the image.



Surfaces

- Light interacts differently with various different objects in the natural world.
 - Lambertian (Random Scattering)
 - Metal (Laws of Reflection)
 - Glass (Snell's Law- Refraction, Total Internal Reflection)
- Each time the ray “hits” an object, a scattered (child) ray is generated.
- Color of the pixel is a combination of the colors observed by all of the generated rays.

Software Implementation



Generate Sub-Rays

- Used for anti-aliasing.
- Generate 100 rays from a single base ray with small random offsets.
- Color of a pixel is cumulative sum of colors seen by all sub-rays.

- Iterates through all objects in scene.
- Solve the quadratic equation to calculate points of intersection.
- Determines the closest valid intersection point.
- Calculate surface normal.


Hit Function

Scatter

- Material dependent light interaction-
- Lambertian spheres- Scattered direction is a random offset from normal direction.
- Metallic Spheres- Scattered direction is determined by law of reflection.

SW Baseline

- Baseline SW code runs all functions on the arm processor.
- Time-
 - 49.36 s (Float)



Pre-Implementation: Profiling

Flat profile:

Each sample counts as 0.01 seconds.

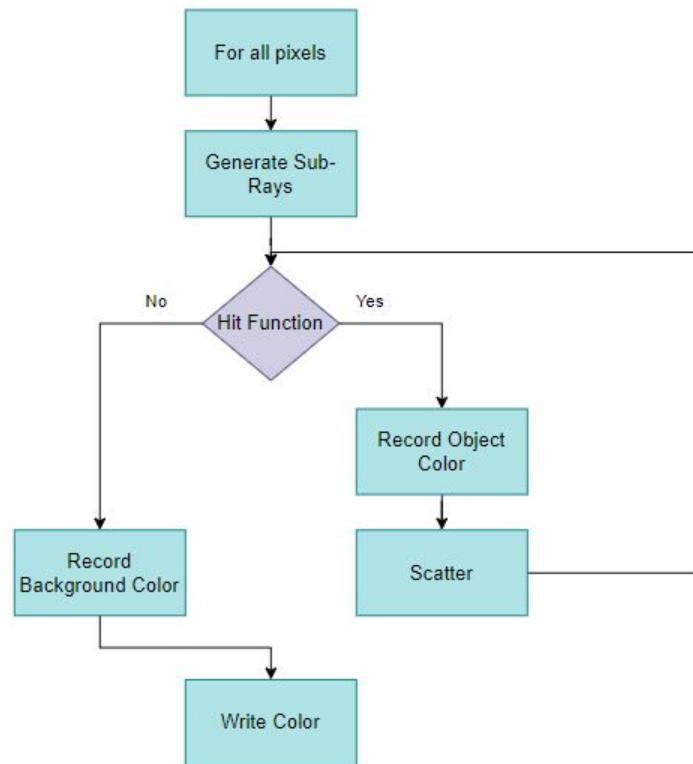
% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
25.35	17.94	<u>17.94</u>	1035825045	0.00	0.00	sphere::hit(ray const&, double, double)
16.98	29.96	12.02	2090382266	0.00	0.00	vec3::length_squared() const
8.32	35.85	5.89	2083131383	0.00	0.00	ray::direction() const
7.84	41.40	5.55	7143621	0.00	0.01	hittable_list::hit(ray const&, double, c
7.20	46.49	5.10	1049107284	0.00	0.00	operator-(vec3 const&, vec3 const&)
6.61	51.17	4.68	1044893347	0.00	0.00	dot(vec3 const&, vec3 const&)
3.70	53.79	2.62	1141339453	0.00	0.00	vec3::vec3(double, double, double)

Pre-Implementation: Quantization

- Hit function in SW - Converted from floating point to fixed point.
- Scale Factor of the form 2^N .
- Downscaled the intermediate values dynamically to prevent overflow.
- 49.1s (Float) \rightarrow 71.21 s(Quantized)

Implementation-1


- Implemented a Quadratic Equation Solver in Verilog (Single Object).
- Running at 250 MHz.
- Time- 145.14s



Vivado Block Diagram

Types of Headache

Migraine




Hypertension




Stress



Xilinx Vivado



RTL



HLS



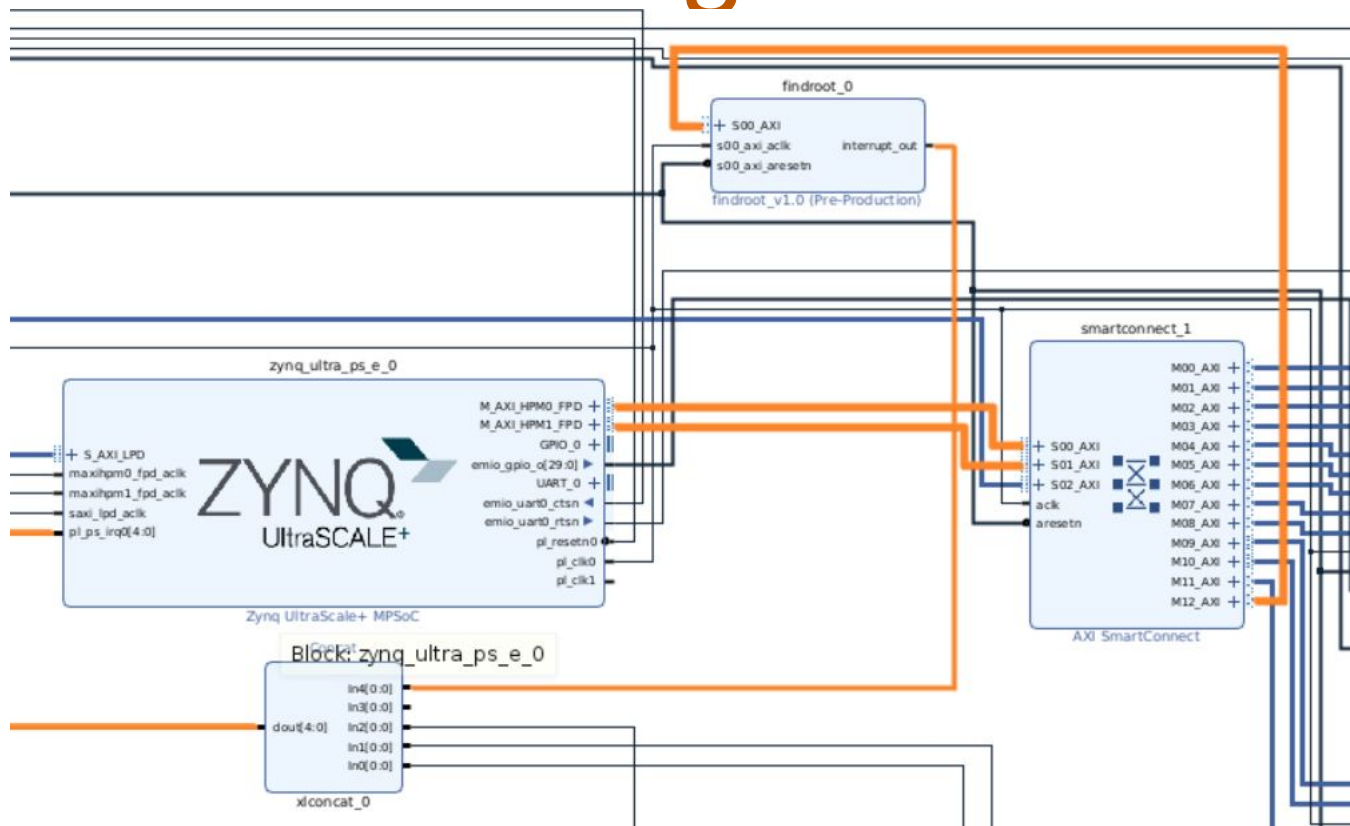
Schematic



Manually
writing the
entire
bitstream using
Notepad.



Vivado Block Diagram



FSM HW design

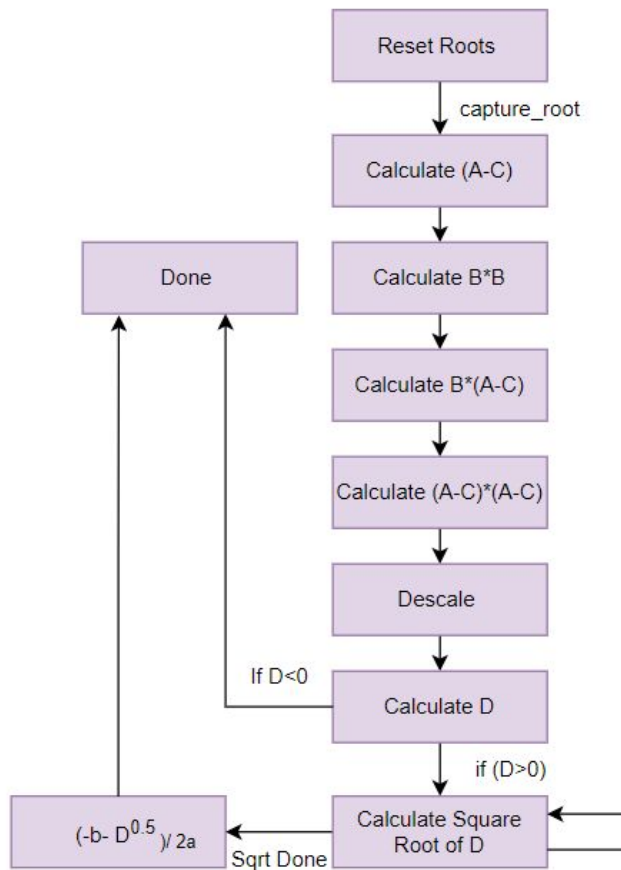
- Essentially calculates solution for-

$$t^2 * (B \cdot B) + 2t * ((B) \cdot (A - C)) + (A - C) \cdot (A - C)$$

$$(-b - D^{0.5}) / 2a$$

$$\text{Where } D = b^2 - 4a.c$$

- Simplified FSM (Takes about 45 cycles (15+30) in the worst case).



Square Root Calculation

- C code provided by Dr.McDermott as part of Lab2
 - Converted C algorithm to Verilog.
- Takes about 15/30 cycles to calculate square root.

```

unsigned long int_sqrt(int n)
{
    int root = 0;
    int bit;
    int trial;

    bit = (n >= 0x10000) ? 1<<30 : 1<<14;
    do
    {
        trial = root+bit;
        if (n >= trial)
        {
            n -= trial;
            root = trial+bit;
        }
        root >>= 1;
        bit >>= 2;
    } while (bit);
    return root;
}
    
```

```

module SquareRoot(
    input wire [31:0] n,
    input wire clk,
    input wire capture,
    output reg[31:0] root,
    output wire done
);
    reg [31:0] bit;
    wire [31:0] trial;
    wire [31:0] subtraction;
    wire [31:0] addition;
    reg [31:0] num;

    assign trial = root + bit;
    assign done = (bit==0);

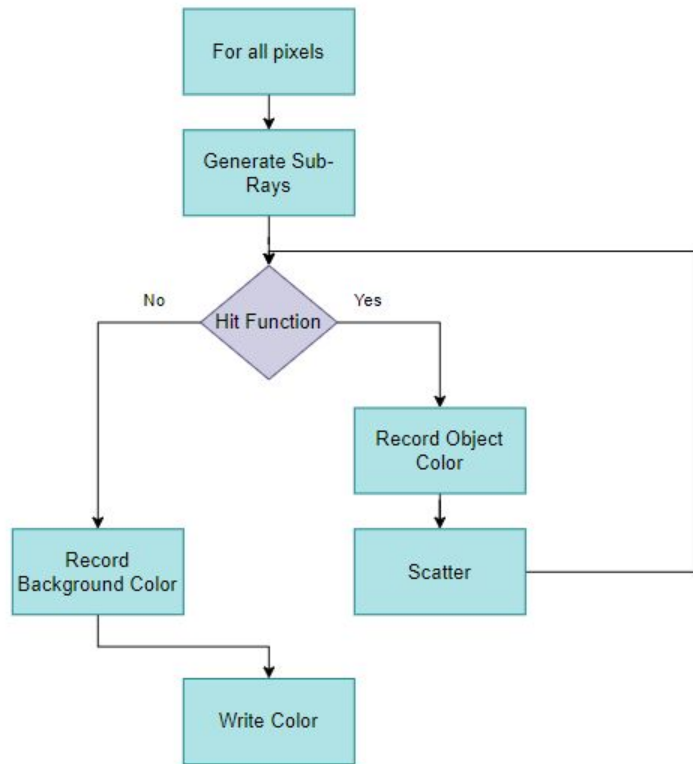
    always @(posedge clk) begin
        if(capture == 1'b1)
            begin
                num<=n;
                root<=0;
                bit<=0;
                if(n==32'h00010000)
                    bit<=1<<30;
                else
                    bit<=1<<14;
            end
        else if(bit != 0) begin
            if(num>=trial) begin
                num<= num - trial;
                root<=(trial + bit)>>1;
                bit<= bit>>2;
            end
            else begin
                root<= root>>1;
                bit<= bit>>2;
            end
        end
    end
end
    
```

Application Code

- Send scene, ray origin and direction information to HW.
- Detect whether a ray hits a particular object in HW. This is done iteratively for all objects and all rays in SW.
 - #Transfers/HW call = 11 integer inputs , 2 integer outputs
- Communicated with slave registers directly using mmap.
- Used Interrupts to detect completion of root calculation.
 - #Interrupts = $400 * 225 * 100 * 10 * 4 = 3.6e8$ (worst case)!

Implementation-2

- Implemented a Quadratic Equation Solver in HLS (Single Object).
- Time- 634.02s
 - Note that this was synthesized at 100MHz.

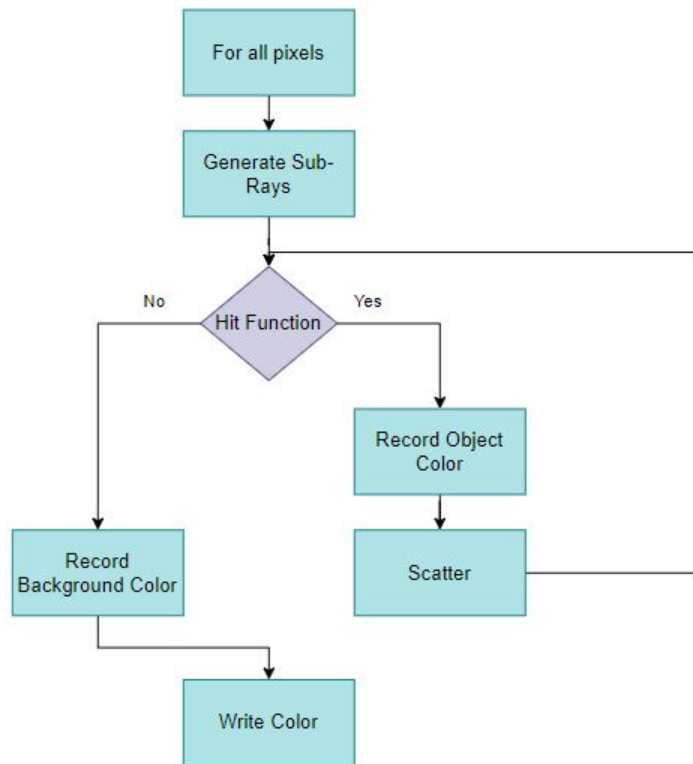


Interlude: High Level Synthesis

- Similar to writing C/C++ Code.
- Can use #pragmas to define AXI ports and perform code optimizations like loop unrolling and pipelining.
- C Simulation (with C/C++ testbench)- Used to verify correctness.
- C Synthesis- To generate verilog code.
- C/RTL Co-simulation- To test the synthesized code with the same C testbench.
- Export RTL- To export the final code as an IP block (for use in Vivado).

Implementation-3

- **Multi-Object** Quadratic Solver.
- Parallelized implementation at 250 MHz.
- Time- 176.73s(polling)



Salient Features

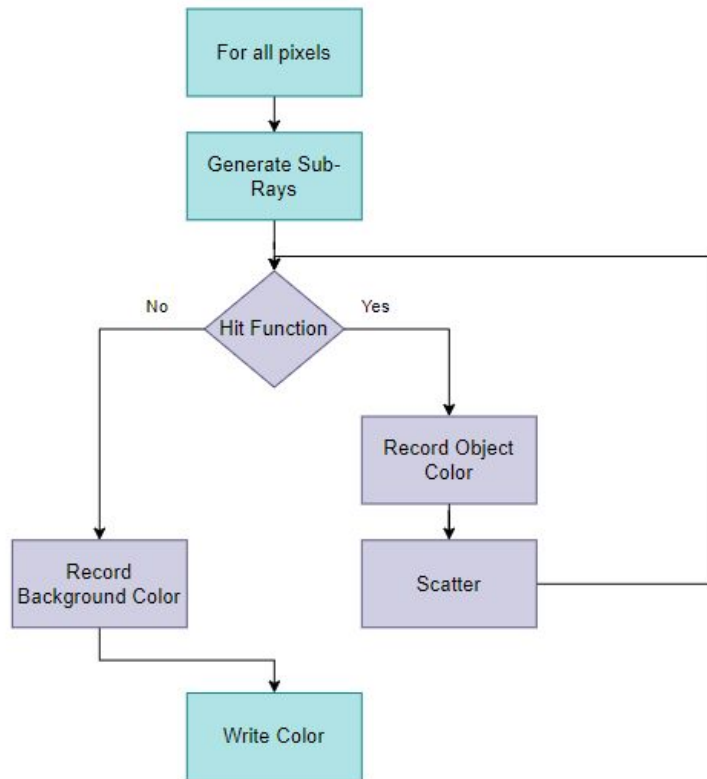
- Vivado design remains unchanged.
- The SW code had to be refactored, which slowed it down.
 - HW implementation still gives overall speedup.
- Reduced #Interrupts by a factor of 4.
- Reduced redundancy in data transfers (scene and geometry information).
- Experimented with polling - Since processor has no work.

Implementation-3 Time Summary


- Overall reducing the number of data transfers and interrupts led to a time of 273.92s
- Changing to polling reduced it to 176.73s

Implementation-4

- Offloaded Ray Scattering operation to HW.
- Time- 35.42s



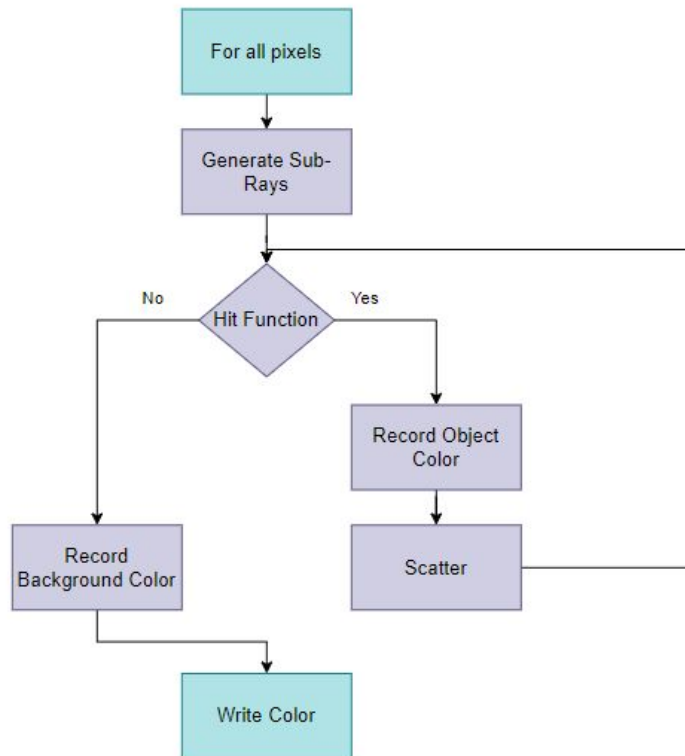
Salient Features



- Vivado design remains unchanged.
- Implemented a 32-bit LFSR to emulate random scattering in Lambertian surfaces.
- Represented colors using fixed point integers.
- Reduced #Interrupts by a factor of 10.
- Further reduced data transfers (ray origin information).

Implementation-5

- Generated sub-rays (100) in HW.
- Leveraged previously developed LFSR to generate sub-rays.
- Reduced #Interrupts by a factor of 100.
- Time- 9.72s

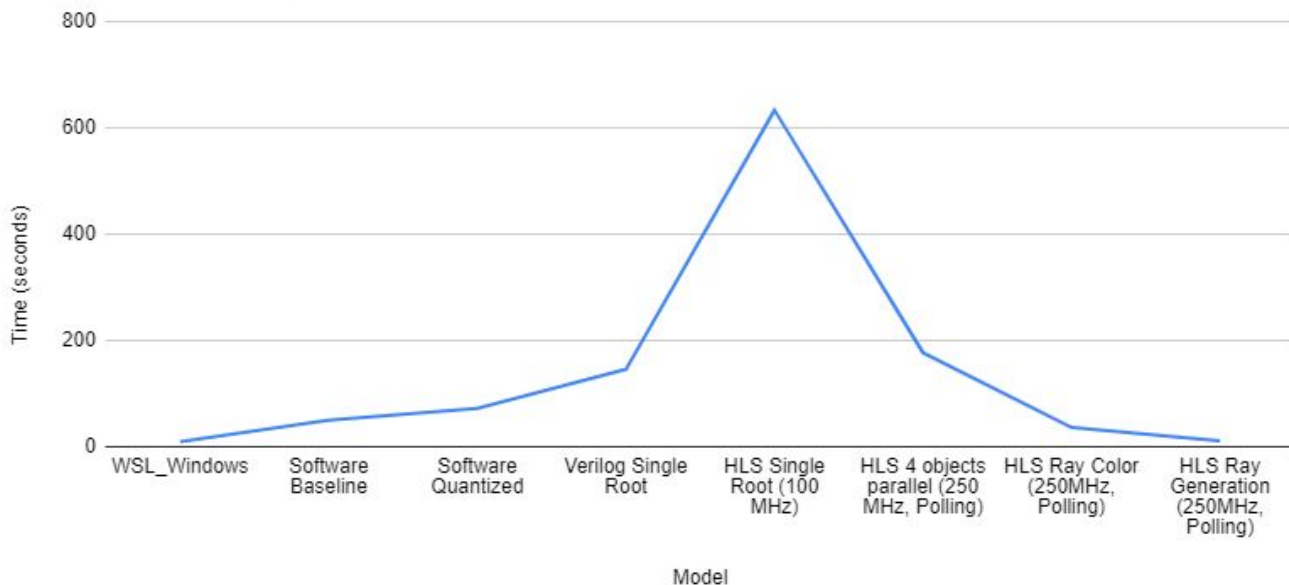


Performance Comparison

Model	Time (seconds)
WSL_Windows	8.64
Software Baseline	49.36
Software Quantized	71.21
Verilog Single Root	145.14
HLS Single Root (100 MHz)	634.023
HLS 4 objects parallel (250 MHz, Polling)	176.73
HLS Ray Color (250MHz, Polling)	35.42
HLS Ray Generation (Final_final_actual_final)	9.72

Performance Graph

Performance Comparison for all models



Debug Techniques

- Verilog Debugging Techniques:
 - Used File Handling to verify functionality for a large number of test cases.
 - Registers for debug
- Modified HLS Files into a C++ Header File for functional verification.
- Visual Debugging
- GDB

Scope for Future Work

- Expand to cover more shapes/materials.
- Scale algorithm for more number of objects (more parallelization).
- Implementing techniques like BVH in hardware.
- Explore Texture Mapping.
- Parallelize the verilog code.

References

- <https://raytracing.github.io/books/RayTracingInOneWeekend.html>
- <https://www.scratchapixel.com/index.html>
- <https://github.com/ssloy/tinyrenderer/wiki>
- These are wonderful references to understand how graphics works from first principles, please do have a look at it!

Thank Yyou

- Prof. Mark McDermott
- TA- Abhijith Venkkateshraj

Demo

